

# Design and compilation of a C-like front end language for GPRM

Ross Meikleham

27th March 2015

# Introduction


- In the past most code written with serial computation in mind
- Rate of increase of processor speeds declining, focus on multiple processor cores
- Would be useful to be able to reuse existing serial code to take advantage of multiple cores without rewriting it.

# GPRM

- The Glasgow Parallel Reduction Machine(GPRM) is a framework for parallel programming using a task based approach.
- Seperate code into task code and communication code.
- Task code is written as C++ classes, can use already written serial C++ code
- Communication code written in a language called GPIR

# GPRM Structure

**GPRM is called from C++ code, the GPRM may return a value back to the calling C++ code**



**GPIR code is compiled down to GPRM Bytecode to be run on the GPRM.**

**GPIR code manages the composition of Kernel Class methods and threads to run the methods on.**

**Task Kernels must reside within the GPRM::Kernel Namespace, these are written in C++.**

**Kernel methods may return values back to the GPRM**

# GPIR

- Glasgow Parallel Intermediate Representation language
- Purely functional S-expression like language like lisp or scheme.
- Evaluated in parallel by default with optional sequential evaluation
- Controls the composition of tasks.

# Problems with GPIR

- Language isn't “consistent” with the rest of the framework.
- Requires programmers to learn multiple languages which are vastly different.
- Still requires programmers to manually enter threads for each individual task to run on

# GPC

- GPC (Glasgow Parallel C). Near subset of C++ with restrictions.
- Layer on top of GPIR, can take advantage of parallel evaluation semantics of GPIR.
- Introduces two new keywords “seq” and “par” to denote sequential or parallel evaluation of code blocks.

# GPC Features

- Purely functional, no multiple assignments
- Support for C++ objects, basic arithmetic operations, functions, types, restricted for-loops.
- Serial semantics.
- Jumping to labels is an expensive operation in the GPIR so inlining as much as possible to the point that execution path can be determined at compile time is useful.
- However kernel calls for C++ can return values, for a turing complete language this is essentially the halting problem.



# Kernel type system

- Resulting return values from calls to C++ Kernel Object methods placed in the `GPRM::Kernel` namespace.
- Kernel types can be passed into kernel calls, but not used as “GPRM” types, this prevents them being used for conditional statements, caught by type checker.
- Essentially keeps GPC purely functional while the “side effect” code is restricted to the kernels.

# Benefits of being a C++ subset

- Being a subset of C++, we can define `seq` and `par` as no-ops, compiler ignores preprocessor directives.
- Can be compiled using a C++ compiler, due to serial semantics behaviour should be the same as parallel version.
- Can use already available C++ tools to test.

# GPC Implementation

- Language compiler is implemented in Haskell
- Uses parsec library to parse/lex source
- Type and Scope checking then run over the AST
- Then interprets the AST to generate GPIR code

# Interpreter

- Runs over AST
- Assigns threads to tasks
- Performs branch elimination on if statements.
- Constant propogation and partial evaluation of expressions
- Function inlining for speed and for loop unrolling
- Generates GPIR AST, can be “pretty printed” to GPIR source

# MergeSort Example

```
GPRM::Kernel::Mergesort MS;

int overall_size = 8;
int cutoff = 1;

void Sort(int size, int c, int index) {
    int new_size = size/2;
    if (size <= cutoff) {
        MS.serial_ms(c%index, size);
    } else {
        MS.merge_two(c%index, size,
            Sort(new_size, 2*c, 2*index),
            Sort(new_size, 2*c + 1, 2*index)
        );
    }
}

void GPRM_MergeSort() {
    seq {
        MS.array(overall_size);
        Sort(overall_size, 1, 1);
    }
}
```

# MergeSort Example

```
;GPRM_MergeSort.yml
```

```
(seq
  ' (MS.array[0] '8)
  ' (MS.merge_two[0] '0 '8
    (MS.merge_two[1] '0 '4
      (MS.merge_two[2] '0 '2
        (MS.serial_ms[3] '0 '1)
        (MS.serial_ms[4] '1 '1))
      (MS.merge_two[5] '1 '2
        (MS.serial_ms[6] '2 '1)
        (MS.serial_ms[7] '3 '1)))
    (MS.merge_two[0] '1 '4
      (MS.merge_two[1] '2 '2
        (MS.serial_ms[2] '4 '1)
        (MS.serial_ms[3] '5 '1))
      (MS.merge_two[4] '3 '2
        (MS.serial_ms[5] '6 '1)
        (MS.serial_ms[6] '7 '1))))))
```

# Summary

- Designed a purely functional language which is familiar to programmers which describes the composition of serial C++ tasks to be run in parallel
- Built the compiler for this language.
- Questions?